

Accelerating Convolutional Neural Networks through CUDA Optimizations on GPUs

René Richard Enjilian
Technische Universität Berlin

Introduction

- Goal:** Implement and optimize a CNN from scratch on a GPU using CUDA
- Approach:** Created two CNN implementations on the GPU – a naive (unoptimized) baseline and an optimized using various CUDA techniques
- Importance:** GPU acceleration reduces CNN training and inference times, making deep learning more efficient
- Result:** Demonstrated notable performance improvements through targeted CUDA optimization techniques

Convolutional Neural Network

- Convolutional Neural Networks (CNN) are neural networks specifically designed to handle structured data with spatial dependencies, such as images or videos.
- Applications:** Widely used in image recognition, autonomous driving, medical imaging, and real-time video analysis
- Convolutional Layer**
 - Apply kernels (filters) to input data to detect local spatial patterns:

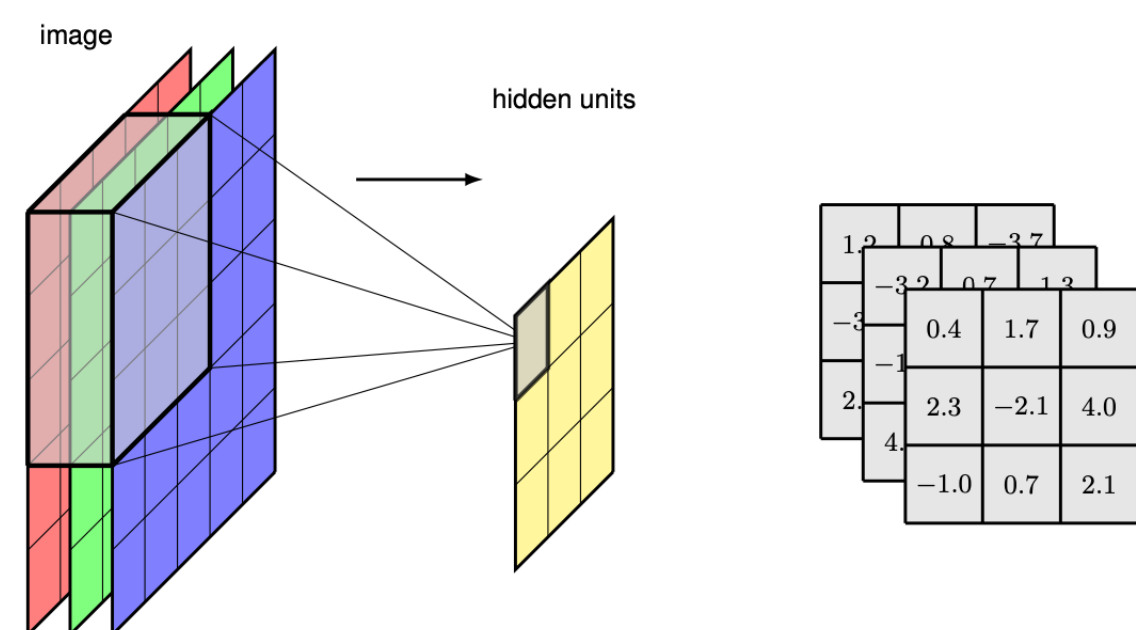


Figure 1: Illustration of the convolution operation. An input image (left) is convolved with a kernel (right), producing a feature map (center).

- Rectified Linear Unit (ReLU)**
 - Non-linear activation function allowing the CNN to model complex, non-linear relationships in the data
 - $ReLU(x) = \max(0, x)$
- Pooling Layer (Max-Pooling)**
 - Reduces spatial dimensions, decreases the number of parameters, and prevents overfitting:

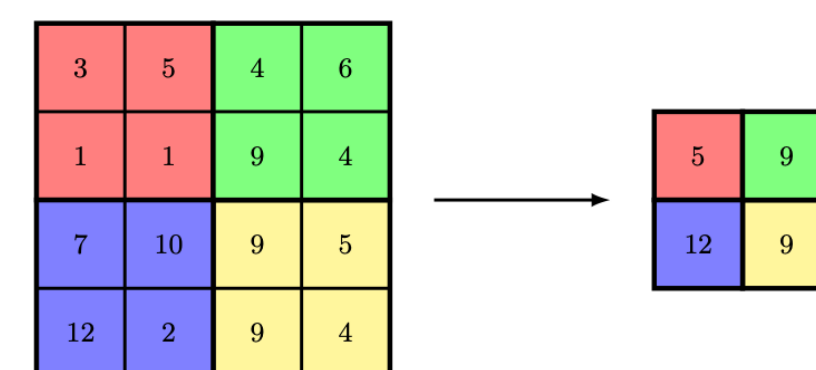


Figure 2: Illustration of max-pooling, reducing each region (coloured) to its maximum value.

- Fully Connected Layer**
 - Combine extracted features to perform final classification
 - Each neuron connects to every output from the previous layer
- Softmax Activation**
 - Converts raw scores into probabilities
 - Indicates the likelihood of each class
 - $Softmax(x_i) = \frac{e^{x_i}}{\sum_j e^{x_j}}$

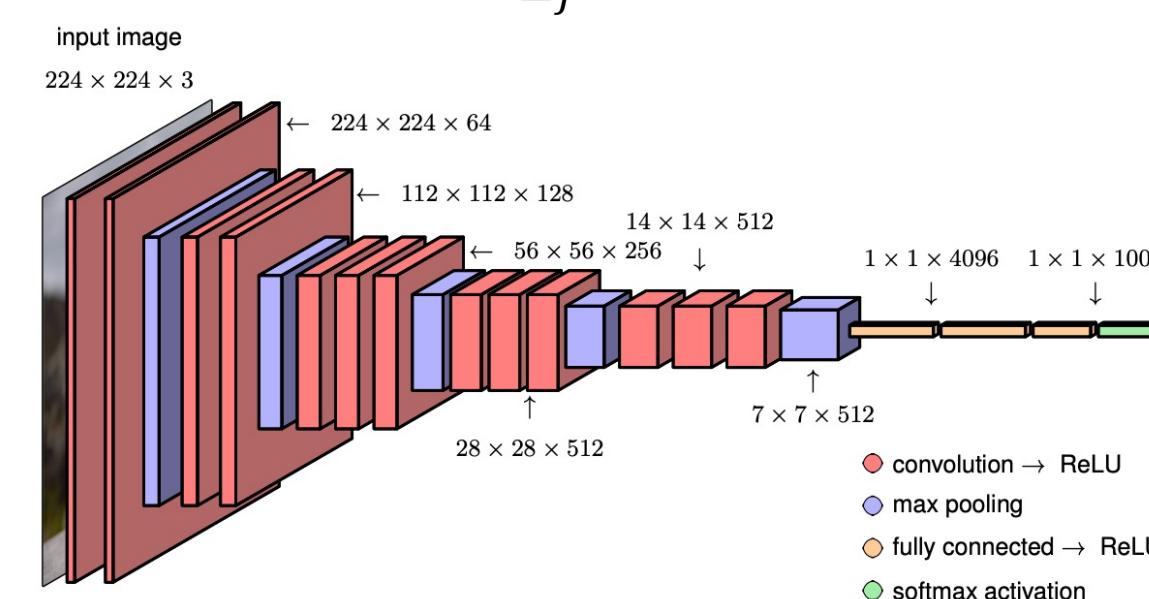


Figure 3: Example CNN architecture (VGG-16).

Implementation & Optimizations

Operation (Kernel)	Baseline (Naive)	Optimized	Performance Gains
Convolution + ReLU Forward	Separate kernels; Direct Global memory loads; No tiling	Kernel Fusion; Shared memory tiling; Constant memory (caching filters)	Fewer kernel launches; Less global mem I/O; Fast filter access
Max-Pooling + Flatten Forward	Separate kernels	Kernel Fusion	Fewer kernel launches; Less global mem I/O
Fully Connected Forward	Naive dot-product per neuron (loop-based)	GEMM-style matrix multiply with shared memory and tiling	Less global mem I/O (tiling); Better parallelization
Softmax + CrossEntropy Forward	Multiple loops for max, exp, sum; Repeated exponent calls	Single-pass approach with local arrays; Fast CUDA intrinsic <code>__expf</code>	Fewer repeated ops (exp, sum); Less overhead from looping
Fully Connected Backward	Two separate kernels: $\nabla W, \nabla b$ and ∇in ; Loops over batch per thread	Shared mem reduction for $\nabla W, \nabla b$; fmaf and loop unrolling for ∇in	Less global mem I/O; Parallel accumulation; Reduced loop overhead
Max-Pooling + Flatten Backward	Separate kernels; Uses atomicAdd	Kernel fusion; Direct index write (no atomic)	Fewer kernel launches; No atomic contention
Convolution Backward	Two separate kernels: $\nabla W, \nabla b$ and ∇in ; Triple nested loops	Warp-level reduction (<code>__shfl_down_sync</code>) for $\nabla W, \nabla b$; Per-sample shared mem for ∇in	Faster gradient reductions; Less global mem I/O
Data Transfer & Streams	Single CUDA stream	Multiple streams; Overlap H2D transfers with kernel execution	Higher GPU utilization; Less idle time on device

Evaluation

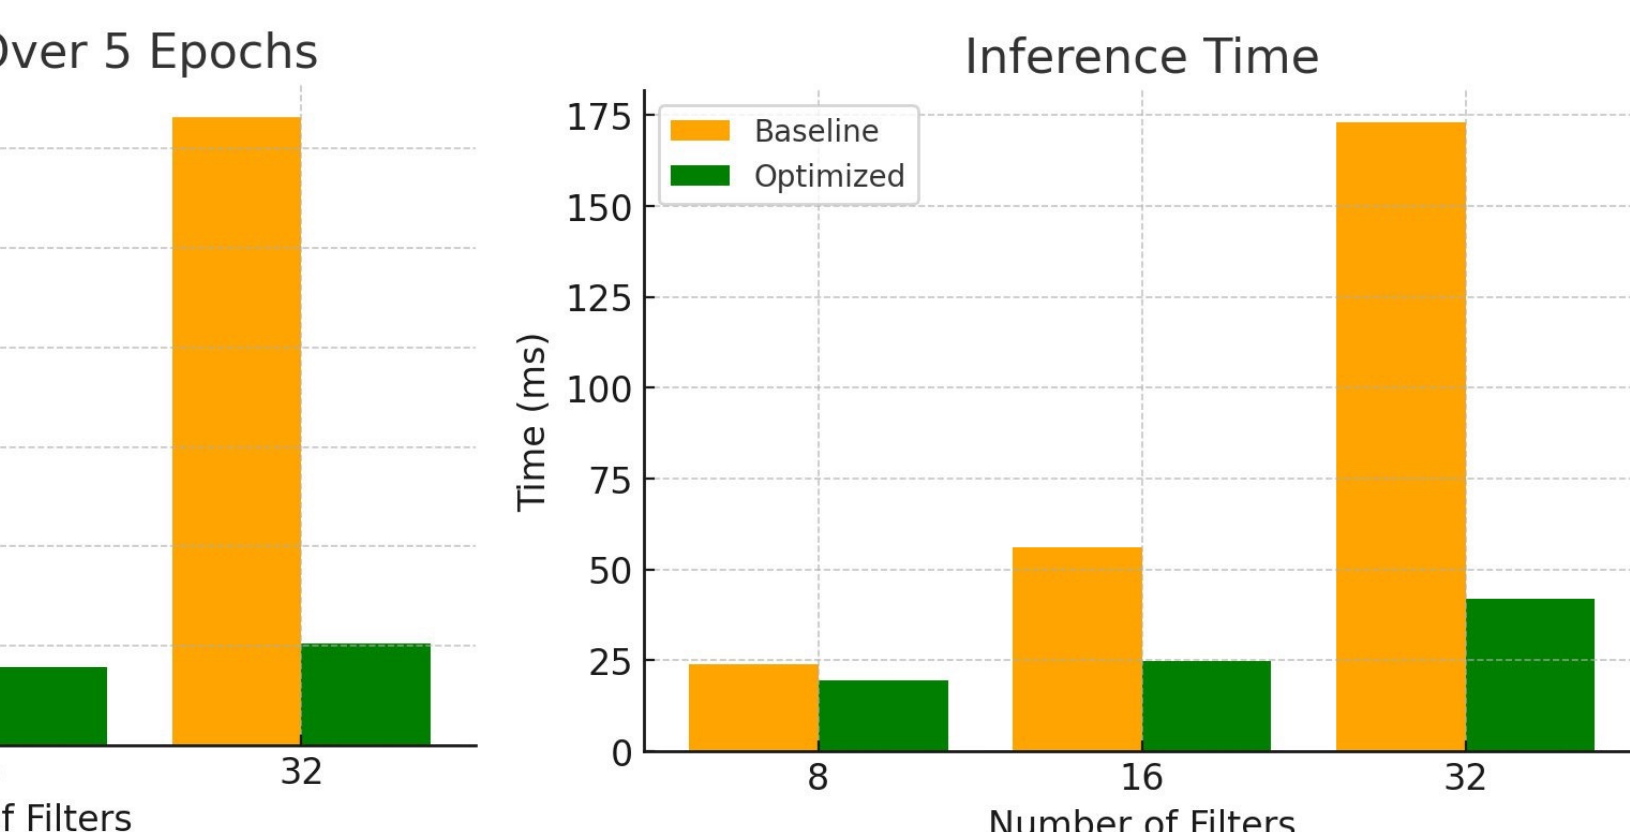


Figure 4: Performance on MNIST dataset.

- Dataset:** MNIST
- Batch size:** 64
- Average Speedup**
 - Training: **6.5x**
 - Inference: **2.5x**

References

C.M. Bishop, *Deep Learning: Foundations and Concepts*, Springer, 2024